

# Application Exercises

*Ken Chen*

*February 24, 2019*

```
setwd("D:/Perspectives/Computational Modeling/hw07")
set.seed(123)
library(tidyverse)
```

```
## -- Attaching packages -----
## v ggplot2 3.1.0      v purrr  0.3.0
## v tibble  2.0.1      v dplyr  0.8.0.1
## v tidyr   0.8.2      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.4.0
```

```
## -- Conflicts -----
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

```
library(readr)
library(ggplot2)
library(rsample)
library(margins)
library(splines)
library(caret)
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## lift
```

```
gss_train = read_csv("./data/gss_train.csv")
```

```
## Parsed with column specification:
```

```
## cols(
```

```
##   .default = col_character(),
```

```
##   age = col_double(),
```

```
##   authoritarianism = col_double(),
```

```
##   childs = col_double(),
```

```
##   con_govt = col_double(),
```

```
##   egalit_scale = col_double(),
```

```
##   income06 = col_double(),
```

```
##   science_quiz = col_double(),
```

```
##   sibs = col_double(),
```

```
##   social_connect = col_double(),
```

```
##   tolerance = col_double(),
```

```
##   tvhours = col_double(),
```

```
##   wordsum = col_double()
```

```
## )
```

```
## See spec(...) for full column specifications.
```

```
gss_test = read_csv("./data/gss_test.csv")

## Parsed with column specification:
## cols(
##   .default = col_character(),
##   age = col_double(),
##   authoritarianism = col_double(),
##   childs = col_double(),
##   con_govt = col_double(),
##   egalit_scale = col_double(),
##   income06 = col_double(),
##   science_quiz = col_double(),
##   sibs = col_double(),
##   social_connect = col_double(),
##   tolerance = col_double(),
##   tvhours = col_double(),
##   wordsum = col_double()
## )
## See spec(...) for full column specifications.
```

## Application Exercises


### Egalitarianism and income

1. Perform polynomial regression to predict `egalit_scale` using `income06`. Use 10-fold cross-validation to select the optimal degree  $d$  for the polynomial based on the MSE. Make a plot of the resulting polynomial fit to the data, and graph the average marginal effect (AME) of `income06` across its potential values. Provide a substantive interpretation of the results.

```
inc06 = select(.data = gss_train, income06)
ega = select(.data = gss_train, egalit_scale)
x_train = cbind(ega, inc06)

mse_lst = rep(0, 16)
cv = vfold_cv(data = x_train, v = 10)
for (i in 1:10){
  splited_set = cv$splits[[i]]
  train = analysis(splited_set); heldout = assessment(splited_set)
  y_true = heldout$egalit_scale
  for (j in 1:16){
    m = glm(egalit_scale ~ poly(income06, j, raw = TRUE), data = train)
    pred = predict(m, newdata = heldout)
    mse = sum((pred - y_true)^2)/length(y_true)
    mse_lst[j] = mse_lst[j] + mse
  }
}

mse_lst = mse_lst/10
tibble_poly = tibble(Training_MSE = mse_lst, Degree = 1:16)
tibble_poly %>%
  ggplot(aes(x = Degree, y = Training_MSE)) + geom_point() + geom_line()
```



As we can see from the plot, the optimal degree is 2.


```
poly_m = lm(egalit_scale ~ income06 + I(income06^2), data = gss_train)
```

```
cplot(poly_m, "income06", what='prediction', draw = F) %>%
  ggplot(aes(x = xvals)) +
  geom_line(aes(y = yvals)) +
  geom_line(aes(y = upper), linetype = 2) +
  geom_line(aes(y = lower), linetype = 2) +
  geom_rug(data = gss_train, aes(x = income06)) +
  labs(title = "Egalitarianism Prediction", x = 'Income', y = "Prediction")
```

##	xvals	yvals	upper	lower
## 1	1	23.11631	25.28371	20.94890
## 2	2	23.15180	25.04001	21.26359
## 3	3	23.15524	24.79232	21.51817
## 4	4	23.12665	24.54195	21.71134
## 5	5	23.06600	24.29036	21.84165
## 6	6	22.97331	24.03899	21.90764
## 7	7	22.84858	23.78870	21.90846
## 8	8	22.69180	23.53894	21.84467
## 9	9	22.50298	23.28678	21.71917
## 10	10	22.28211	23.02670	21.53752
## 11	11	22.02920	22.75132	21.30708
## 12	12	21.74424	22.45297	21.03552
## 13	13	21.42724	22.12502	20.72946
## 14	14	21.07819	21.76262	20.39376


```
## 15    15 20.69710 21.36290 20.03130
## 16    16 20.28396 20.92501 19.64291
## 17    17 19.83878 20.45044 19.22712
## 18    18 19.36155 19.94356 18.77955
## 19    19 18.85228 19.41247 18.29210
## 20    20 18.31096 18.86919 17.75274
```

## Egalitarianism Prediction



```
cplot(poly_m, "income06", what='effect', draw = F) %>%
  ggplot(aes(x = xvals)) +
  geom_line(aes(y = yvals)) +
  geom_line(aes(y = upper), linetype = 2) +
  geom_line(aes(y = lower), linetype = 2) +
  geom_rug(data = gss_train, aes(x = income06)) +
  labs(title = "Average Marginal Effects of Potential Income Values", x = 'Income', y = "MAE")
```

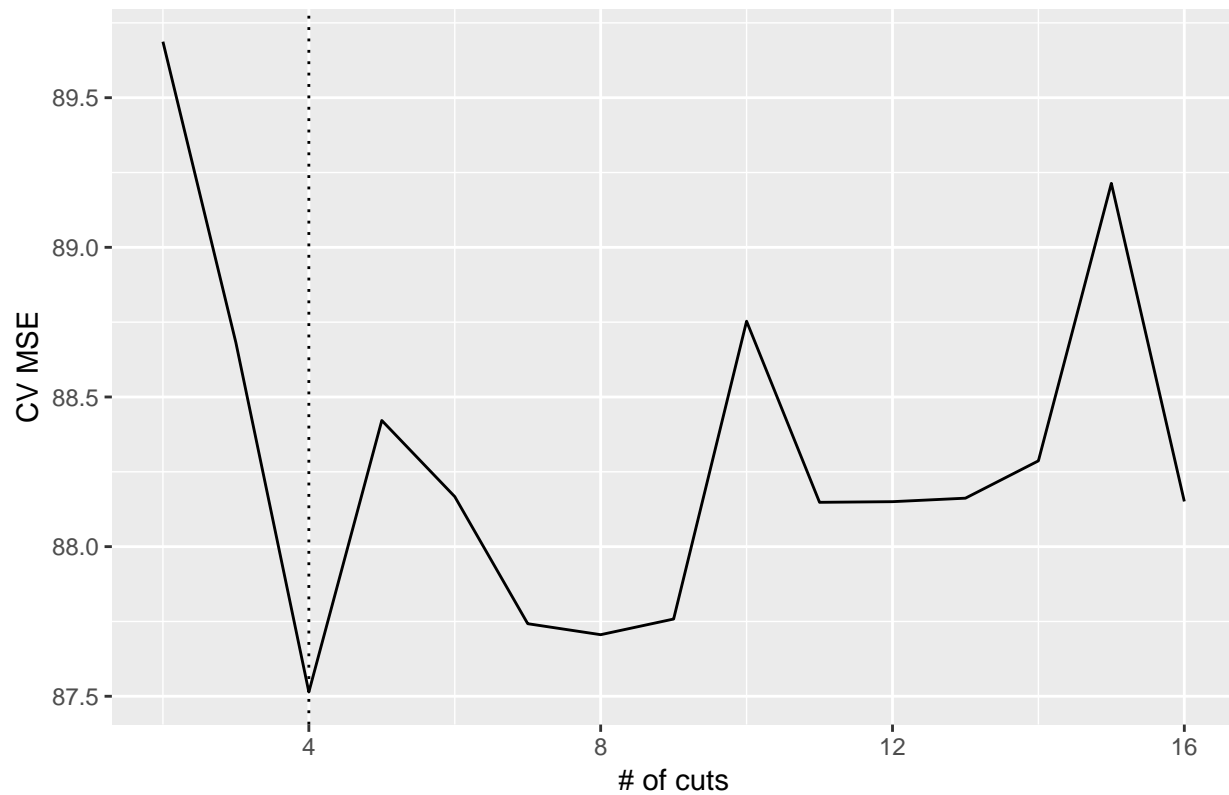
## Average Marginal Effects of Potential Income Values



2. Fit a step function to predict `egalit_scale` using `income06`, and perform 10-fold cross-validation to choose the optimal number of cuts. Make a plot of the fit obtained and interpret the results.

```
mse_lst = rep(0, 15)
for (i in 2:16) {
  gss_train$inc06_cut = cut_interval(gss_train$income06, i)
  m = glm(egalit_scale ~ inc06_cut, data = gss_train)
  mse_lst[i-1] = boot::cv.glm(gss_train, m, K = 10)$delta[1]
}
tibble(cut_num = 2:16, mse = mse_lst) %>%
  ggplot(aes(cut_num, mse)) +
  geom_line() +
  geom_vline(xintercept = which.min(mse_lst) + 1, linetype = 3) +
  labs(title = "Step function regression: Cross-Validation over different cuts", x = "# of cuts", y = "MAE")
```

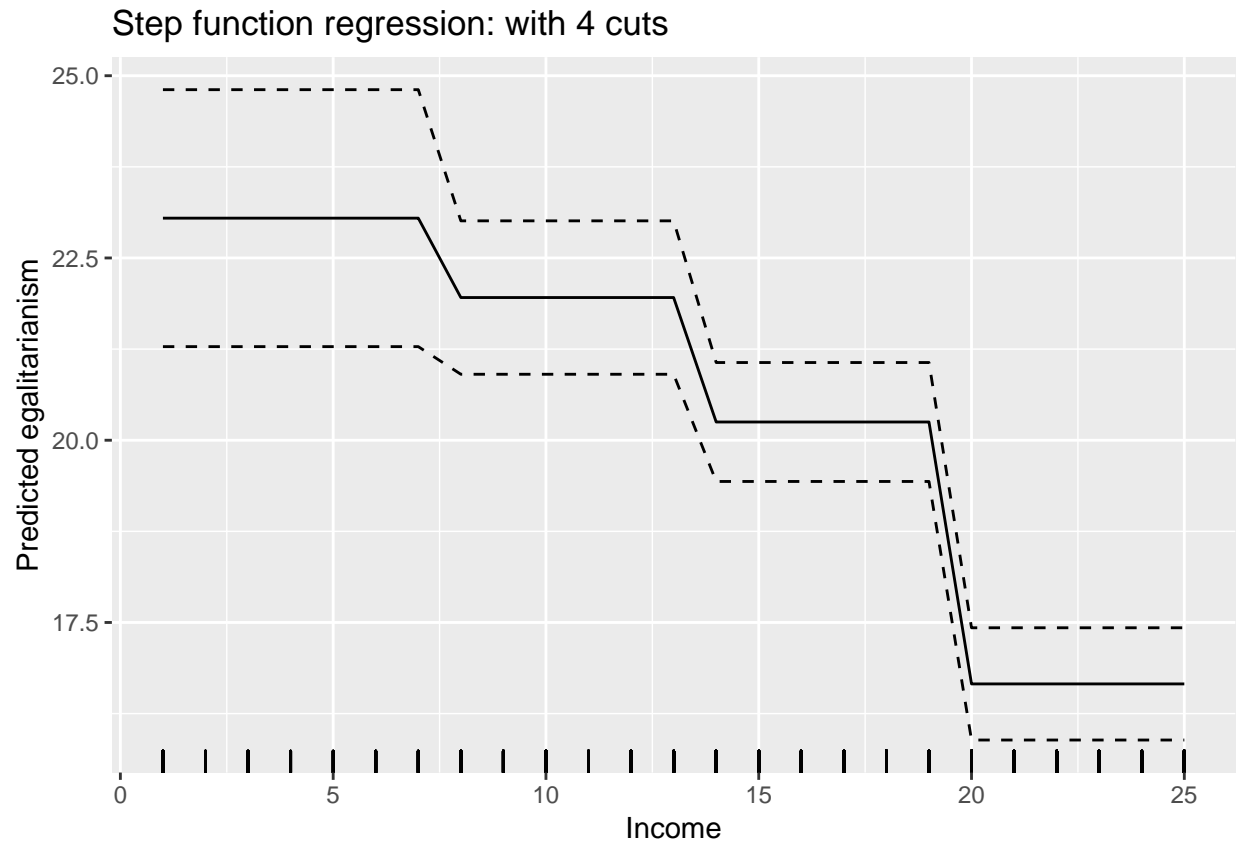
### Step function regression: Cross-Validation over different cuts



The optimal number of cuts is 4 according to our cross validation.

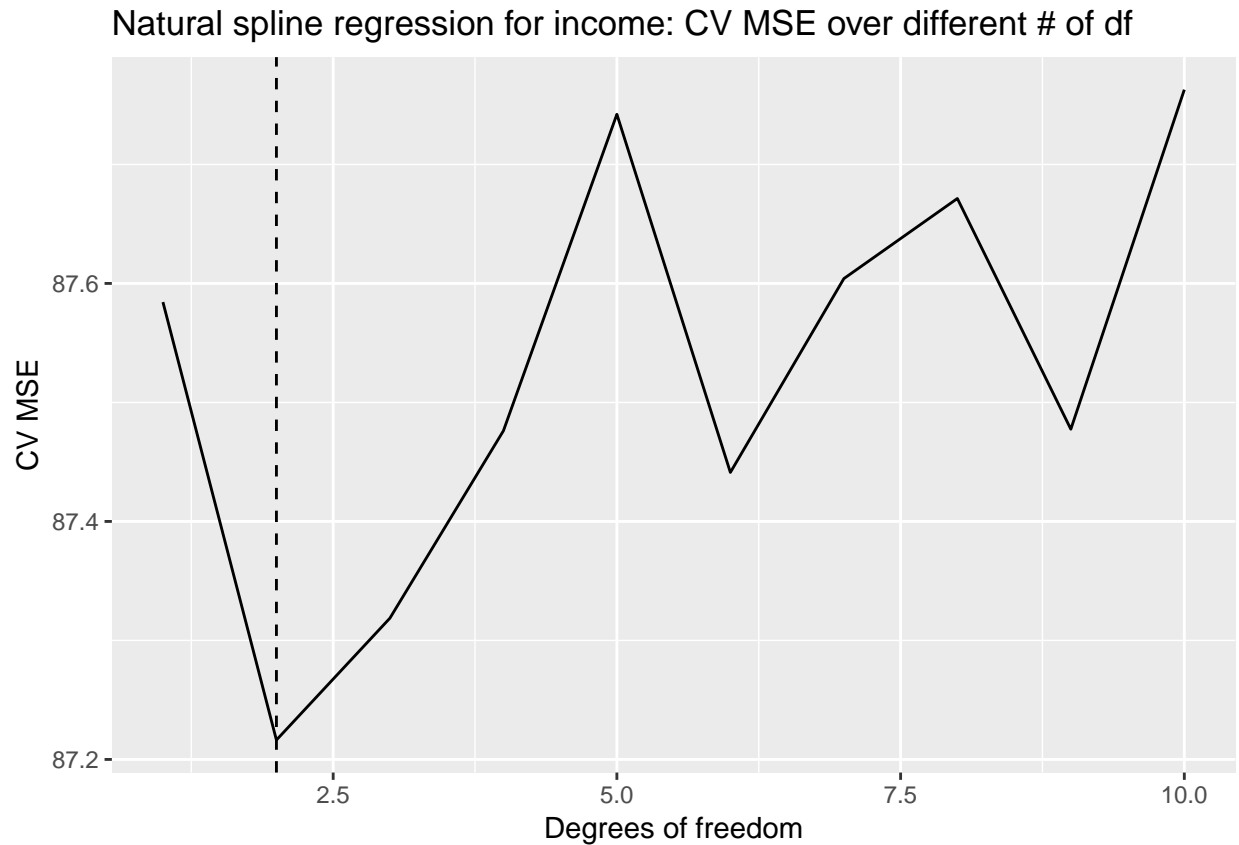
```
m_opt = lm(egalit_scale ~ cut_interval(income06, 4), data = gss_train)
```

```
m_opt %>% prediction %>%  
  ggplot(aes(x = income06)) +  
  geom_line(aes(y = fitted)) +  
  geom_line(aes(y = fitted + 1.96 * se.fitted), linetype = 2) +  
  geom_line(aes(y = fitted - 1.96 * se.fitted), linetype = 2) +  
  geom_rug(data = gss_train, aes(x = income06)) +  
  labs(title = "Step function regression: with 4 cuts", x = "Income", y = "Predicted egalitarianism")
```



3. Fit a natural regression spline to predict `egalit_scale` using `income06`. Use 10-fold cross-validation to select the optimal number of degrees of freedom, and present the results of the optimal model.

```
mse_lst = rep(0, 10)
for (i in 1:10) {
  m_spline = glm(egalit_scale ~ ns(income06, df = i), data = gss_train)
  mse_lst[i] = boot::cv.glm(gss_train, m_spline, K = 10)$delta[1]
}
tibble(df = 1:10, mse = mse_lst) %>%
  ggplot(aes(df, mse)) +
  geom_line() +
  geom_vline(xintercept = which.min(mse_lst), linetype = 2) +
  labs(title = "Natural spline regression for income: CV MSE over different # of df", x = "Degrees of f
```

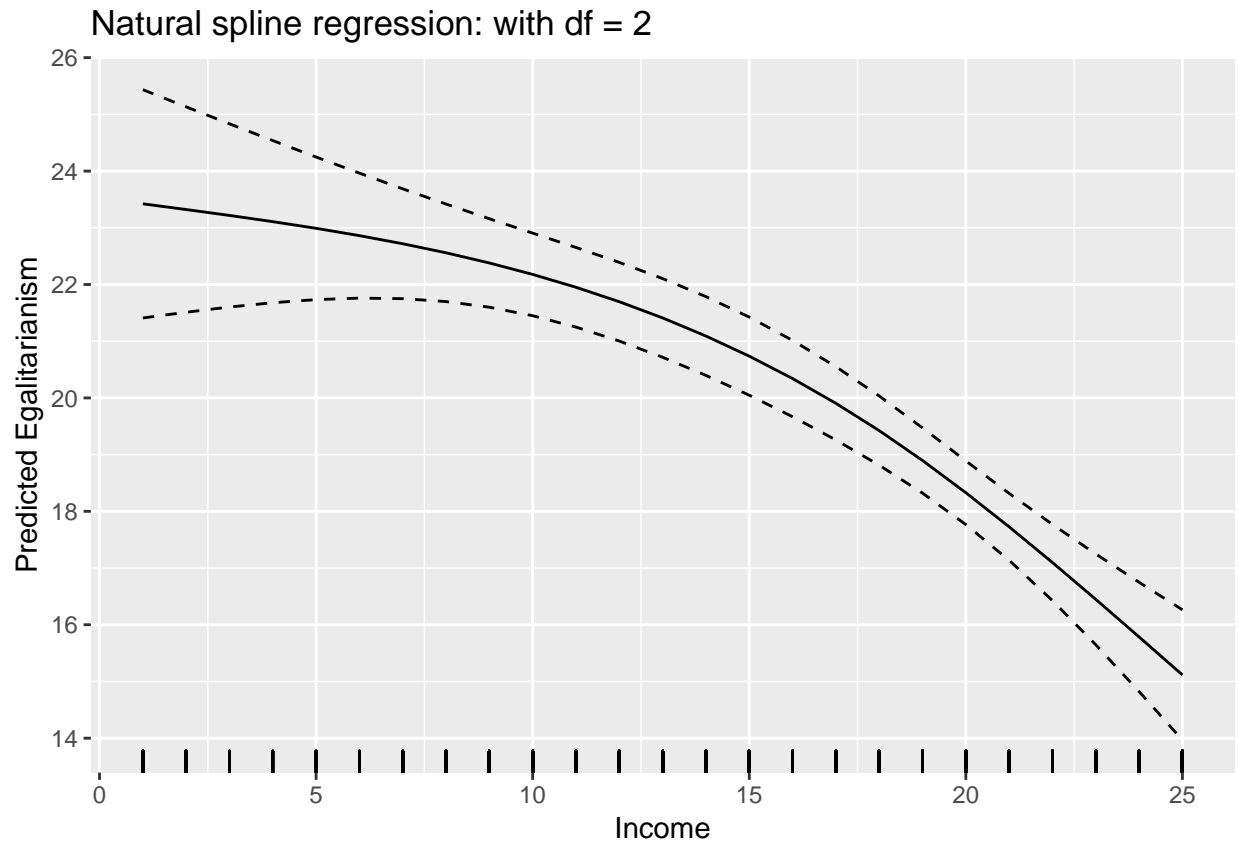


According to the 10-fold cross validation, the optimal degrees of freedom is 2.

```
ns_opt = lm(egalit_scale ~ ns(income06, df = 2), data = gss_train)
```

```
ns_opt %>% prediction %>%
  ggplot(aes(x = income06)) +
  geom_line(aes(y = fitted)) +
  geom_line(aes(y = fitted + 1.96 * se.fitted), linetype = 2) +
  geom_line(aes(y = fitted - 1.96 * se.fitted), linetype = 2) +
  geom_rug(data = gss_train, aes(x = income06)) +
  labs(title = "Natural spline regression: with df = 2", x = "Income", y = "Predicted Egalitarianism")
```






4. Fit a local linear regression model to predict `egalit_scale` using `income06`. Use 10-fold cross-validation to select the optimal bandwidth. Interpret the results.

```
mse_lst = rep(0, 25)
cv = vfold_cv(data = x_train, v = 10)
for (i in 1:10){
  splited_set = cv$splits[[i]]
  train = analysis(splited_set); heldout = assessment(splited_set)
  y_true = heldout$egalit_scale
  j = 1
  for (bdw in seq(0.3, 1.5, 0.05)){
    m = loess(egalit_scale ~ income06, data = train, span = bdw, degree = 1)
    pred = predict(m, newdata = heldout)
    mse = sum((pred - y_true)^2)/length(y_true)
    mse_lst[j] = mse_lst[j] + mse
    j = j+1
  }
}


mse_lst = mse_lst/10
tibble_poly = tibble(Training_MSE = mse_lst, Bandwidth = seq(0.3, 1.5, 0.05))
tibble_poly %>%
  ggplot(aes(x = Bandwidth, y = Training_MSE)) +
  geom_point() +
  geom_line()
```



The optimal bandwidth is 0.85.

```
ggplot(gss_train, aes(income06, egalit_scale)) +
  geom_smooth(method = "loess", span = 0.85, method.args = list(degree = 1)) +
  labs(title = "Local linear regression: with bandwidth = 0.85", x = "Income", y = "Predicted Egalitarianism")
```

Local linear regression: with bandwidth = 0.85



5. Fit a local polynomial regression model to predict `egalit_scale` using `income06`. Use 10-fold crossvalidation to select the optimal bandwidth. Interpret the results.

```
mse_lst = rep(0, 20)
cv = vfold_cv(data = x_train, v = 10)
for (i in 1:10){
  splited_set = cv$splits[[i]]
  train = analysis(splited_set); heldout = assessment(splited_set)
  y_true = heldout$egalit_scale
  j = 1
  for (bdw in seq(0.25, 5, 0.25)){
    m = loess(egalit_scale ~ income06, data = train, span = bdw, degree = 2)
    pred = predict(m, newdata = heldout)
    mse = sum((pred - y_true)^2)/length(y_true)
    mse_lst[j] = mse_lst[j] + mse
    j = j+1
  }
}
```

```
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : pseudoinverse used at 20
```


```
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : neighborhood radius 1
```

```
## Warning in simpleLoess(y, x, w, span, degree = degree, parametric =
## parametric, : reciprocal condition number 0
```

```

mse_lst = mse_lst/10
tibble_poly = tibble(Training_MSE = mse_lst, Bandwidth = seq(0.25, 5, 0.25))
tibble_poly %>%
  ggplot(aes(x = Bandwidth, y = Training_MSE)) +
  geom_point() +
  geom_line()

```




The optimal bandwidth is 5.

```

ggplot(gss_train, aes(income06, egalit_scale)) +
  geom_smooth(method = "loess", span = 5, method.args = list(degree = 2)) +
  labs(title = "Local Polynomial Regression: with bandwidth = 5", x = "Income", y = "Predicted Egalitarianism")

```

### Local Polynomial Regression: with bandwidth = 5



### Egalitarianism and everything

```
library(glmnet)
```

```
## Loading required package: Matrix
##
## Attaching package: 'Matrix'
## The following object is masked from 'package:tidyr':
##
##   expand
## Loading required package: foreach
##
## Attaching package: 'foreach'
## The following objects are masked from 'package:purrr':
##
##   accumulate, when
## Loaded glmnet 2.0-16
```

```
library(pls)
```

```
##
## Attaching package: 'pls'
## The following object is masked from 'package:caret':
```

```
##
##      R2
## The following object is masked from 'package:stats':
##
##      loadings
```

```
library(earth)
```

```
## Loading required package: plotmo
## Loading required package: plotrix
## Loading required package: TeachingDemos
```

```
library(iml)
```

1. Estimate the following models using all the available predictors:
  - a. Linear regression
  - b. Elastic net regression
  - c. Principal component regression
  - d. Partial least squares regression
  - e. Multivariate adaptive regression splines (MARS)
- Perform appropriate data pre-processing (e.g. standardization) and hyperparameter tuning (e.g. lambda for PCR/PLS, lambda and alpha for elastic net, degree of interactions and number of retained terms for MARS)
- Use 10-fold cross-validation for each model to estimate the model's performance using MSE.

```
gss_train = select(gss_train, -inc06_cut)
# Linear Regression Model
lr <- train(egalit_scale ~ ., data = gss_train,
  method = "lm", metric = "RMSE", trControl = trainControl(method = "cv", number = 10), preProcess = c(
)
# Elastic Net Regression
ela.net <- train(egalit_scale ~ ., data = gss_train, method = "glmnet",
  trControl = trainControl(method = "cv", number = 10), metric = "RMSE", preProcess = c("zv", "center",
)
# PCR
pcr <- train(egalit_scale ~ ., data = gss_train, method = "pcr",
  trControl = trainControl(method = "cv", number = 10), metric = "RMSE", preProcess = c("zv", "center",
)
# PLS
pls <- train(egalit_scale ~ ., data = gss_train, method = "pls",
  trControl = trainControl(method = "cv", number = 10), metric = "RMSE", preProcess = c("zv", "center",
)
# MARS
grid <- expand.grid(degree = 1:3, nprune = seq(2, 100, length.out = 10) %>% floor())
mars <- train(egalit_scale ~ ., data = gss_train, method = "earth",
  trControl = trainControl(method = "cv", number = 10), metric = "RMSE", preProcess = c("zv"), tuneGrid
)

summary(resamples(list(
  Linear.Regression = lr,
  Elastic.Net = ela.net,
  PCR = pcr,
  PLS = pls,
```

```

MARS = mars
)))

##
## Call:
## summary.resamples(object = resamples(list(Linear.Regression =
##   lr, Elastic.Net = ela.net, PCR = pcr, PLS = pls, MARS = mars)))
##
## Models: Linear.Regression, Elastic.Net, PCR, PLS, MARS
## Number of resamples: 10
##
## MAE
##
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## Linear.Regression 5.502586 6.014391 6.400598 6.253367 6.566769 6.725281
## Elastic.Net       5.766522 5.949770 6.150565 6.159137 6.385637 6.580547
## PCR               5.899741 6.268928 6.344180 6.449353 6.672020 7.076248
## PLS               5.647323 6.092515 6.282312 6.286030 6.389732 6.999243
## MARS              5.743503 5.914691 6.223570 6.171208 6.461125 6.539609
##
##           NA's
## Linear.Regression    0
## Elastic.Net          0
## PCR                  0
## PLS                  0
## MARS                 0
##
## RMSE
##
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## Linear.Regression 6.789330 7.643416 8.054975 7.898578 8.297012 8.445941
## Elastic.Net       7.023522 7.389858 7.764023 7.728647 8.064425 8.326068
## PCR               7.189889 7.833668 7.996183 8.046503 8.231913 9.077769
## PLS               7.234305 7.790933 7.860970 7.925482 7.932825 8.673940
## MARS              7.378935 7.419016 7.682595 7.781426 8.121300 8.306568
##
##           NA's
## Linear.Regression    0
## Elastic.Net          0
## PCR                  0
## PLS                  0
## MARS                 0
##
## Rsquared
##
##           Min.   1st Qu.   Median     Mean   3rd Qu.
## Linear.Regression 0.2516288 0.2938432 0.3116055 0.3358196 0.3616757
## Elastic.Net       0.2783845 0.2958091 0.3472662 0.3593702 0.4137933
## PCR               0.1776466 0.2797380 0.2955079 0.3066535 0.3529264
## PLS               0.2514042 0.3040409 0.3231183 0.3259740 0.3569333
## MARS              0.2766183 0.3257554 0.3505151 0.3512856 0.3839220
##
##           Max. NA's
## Linear.Regression 0.5056246 0
## Elastic.Net       0.4850429 0
## PCR               0.4442248 0
## PLS               0.4029653 0
## MARS              0.4191424 0

```

Looking at both RMSE and MAE, Elastic Net performed the best among all.

2. Apply model interpretation methods to each model. That is, for each model (the final tuned version), generate permutation-based feature importance plots, PDPs/ICE plots for the five most important variables, and feature interaction plots. Interpret the results with written analysis.

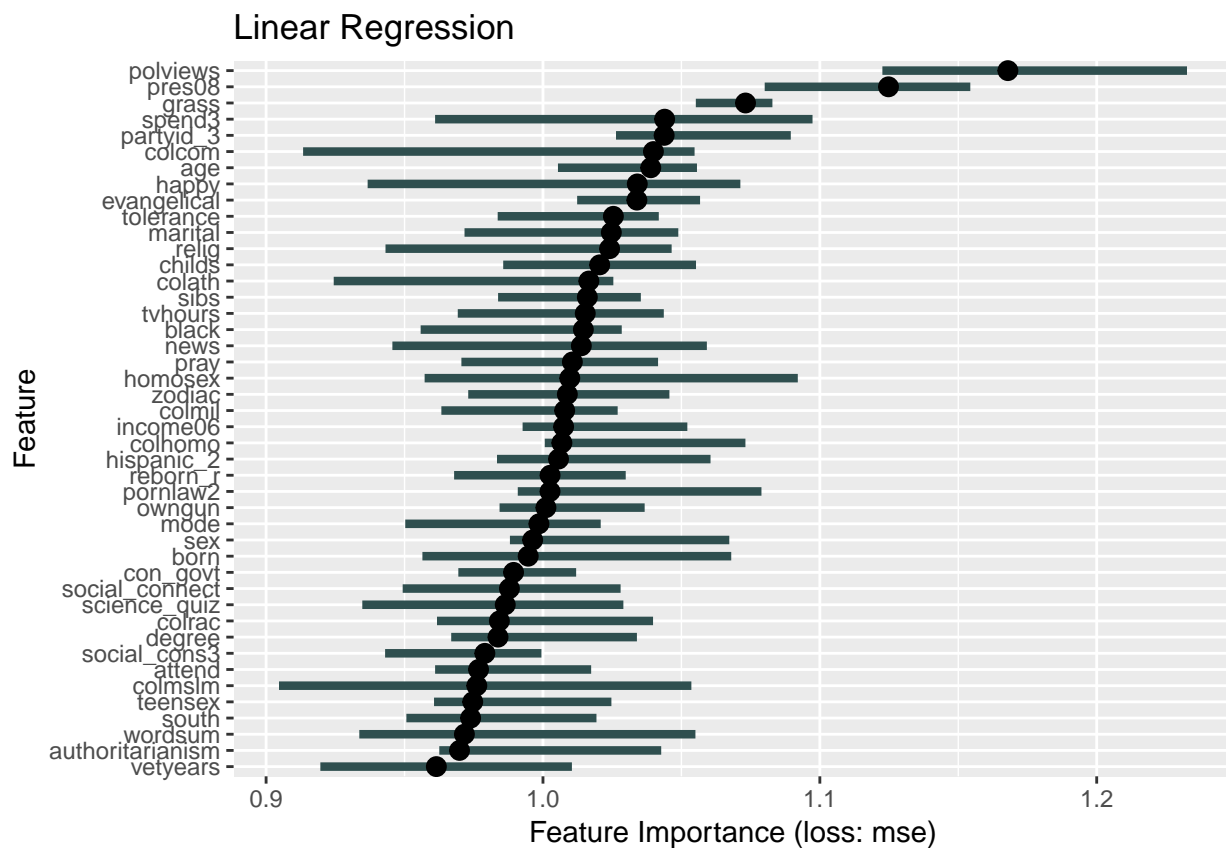
```
pred_lr = Predictor$new( model = lr, data = select(gss_train, -egalit_scale), y = gss_train$egalit_scale)
pred_net = Predictor$new(model = ela.net, data = select(gss_train, -egalit_scale), y = gss_train$egalit_scale)
pred_pcr = Predictor$new(model = pcr, data = select(gss_train, -egalit_scale), y = gss_train$egalit_scale)
pred_pls = Predictor$new(model = pls, data = select(gss_train, -egalit_scale), y = gss_train$egalit_scale)
pred_mars = Predictor$new(model = mars, data = select(gss_train, -egalit_scale), y = gss_train$egalit_scale)
```

```
# Feature Importance
```

```
imp_lr = FeatureImp$new(pred_lr, loss = "mse")
imp_net = FeatureImp$new(pred_net, loss = "mse")
imp_pcr = FeatureImp$new(pred_pcr, loss = "mse")
imp_pls = FeatureImp$new(pred_pls, loss = "mse")
imp_mars = FeatureImp$new(pred_mars, loss = "mse")
```

```
img1 = plot(imp_lr) + ggtitle("Linear Regression")
img2 = plot(imp_pcr) + ggtitle("PCR")
img3 = plot(imp_pls) + ggtitle("PLS")
img4 = plot(imp_net) + ggtitle("Elastic net")
img5 = plot(imp_mars) + ggtitle("MARS")
```

```
img1
```



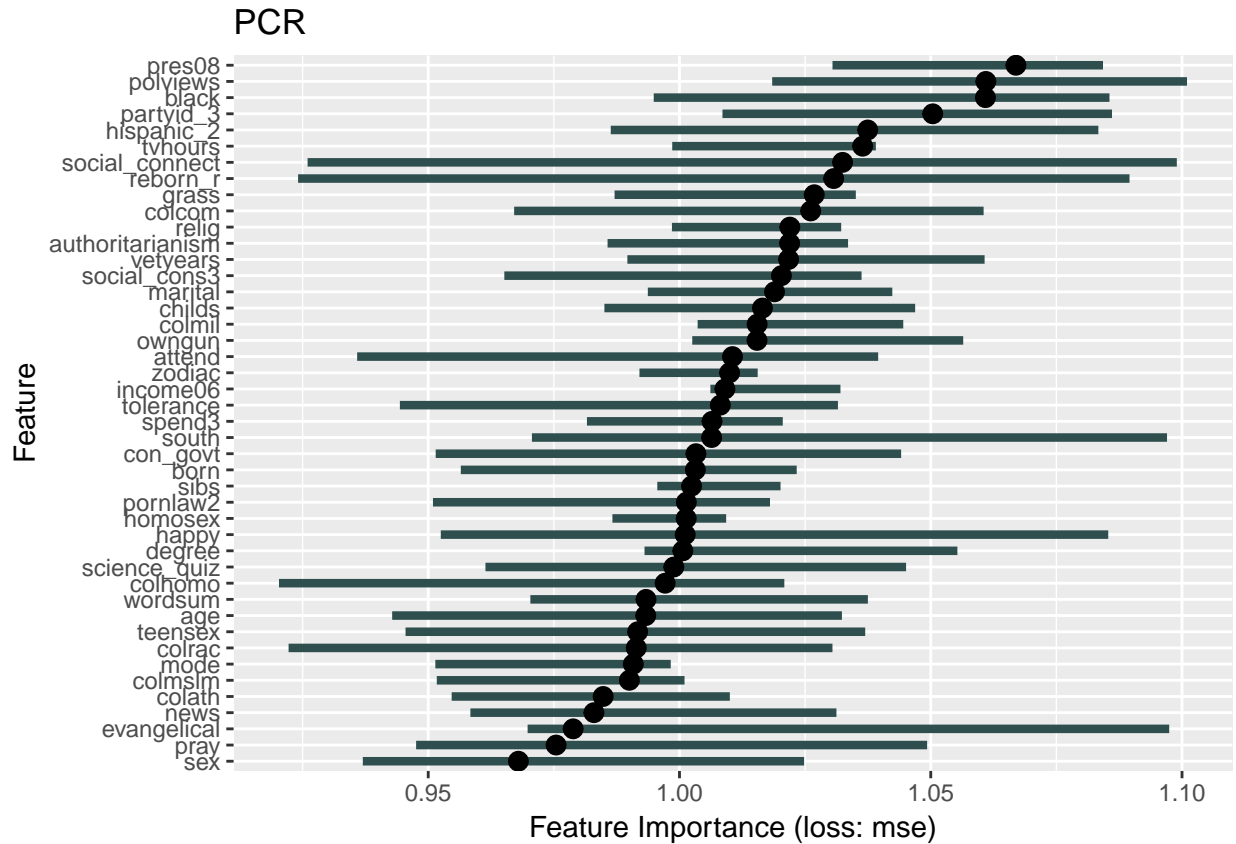
```
head(imp_lr$results, 5)
```

```
##      feature importance.05 importance importance.95 permutation.error
```



```
## 1 polviews      1.1225923  1.167918  1.232604      63.56455
## 2 pres08        1.0800850  1.124796  1.154340      61.21763
## 3 grass         1.0552063  1.073156  1.082864      58.40708
## 4 spend3        0.9610456  1.043940  1.097356      56.81702
## 5 partyid_3     1.0263772  1.043772  1.089503      56.80784
```

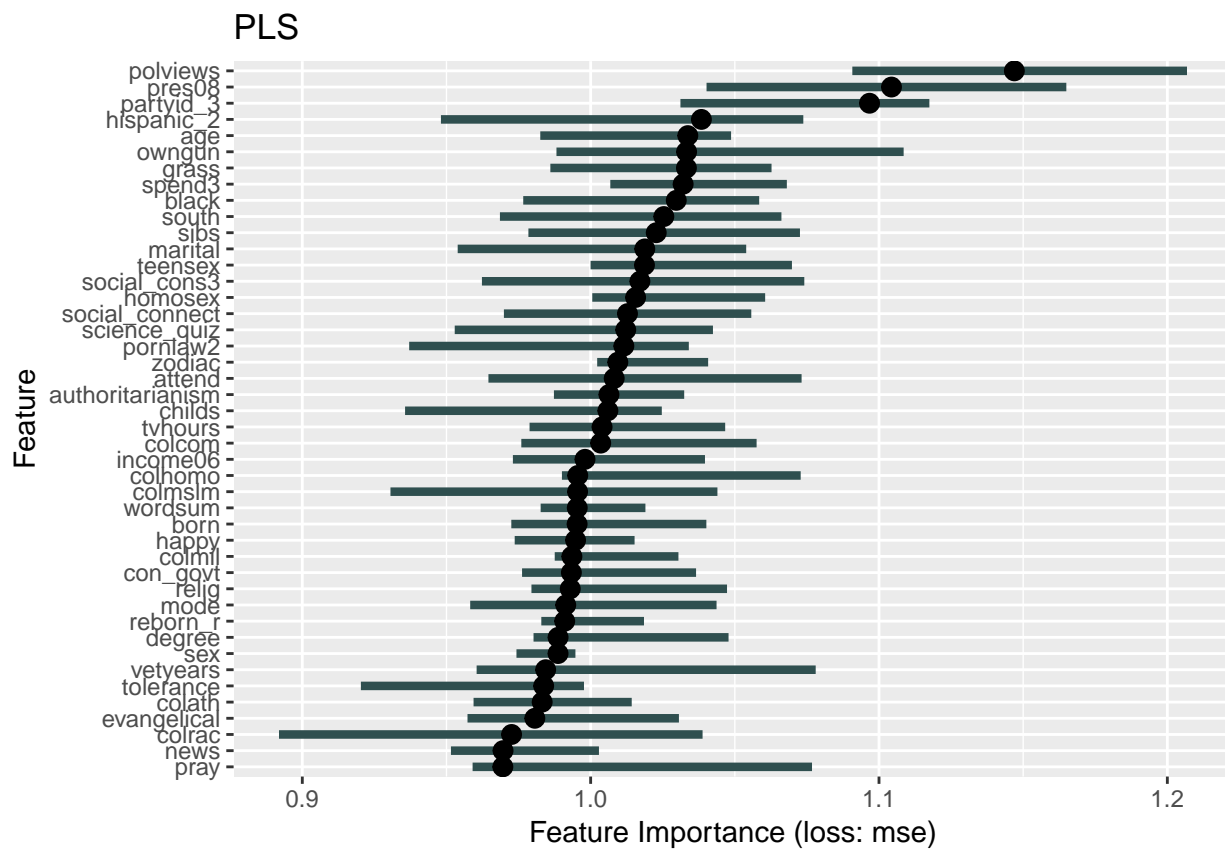
img2



```
head(imp_net$results, 5)
```

```
##      feature importance.05 importance importance.95 permutation.error
## 1 pres08      1.1220595  1.178630  1.205459      67.18219
## 2 polviews    1.0599580  1.134426  1.170556      64.66255
## 3 partyid_3   1.0183967  1.063340  1.087106      60.61060
## 4 degree     0.9799325  1.053507  1.060483      60.05011
## 5 news       0.9377634  1.045942  1.053519      59.61895
```

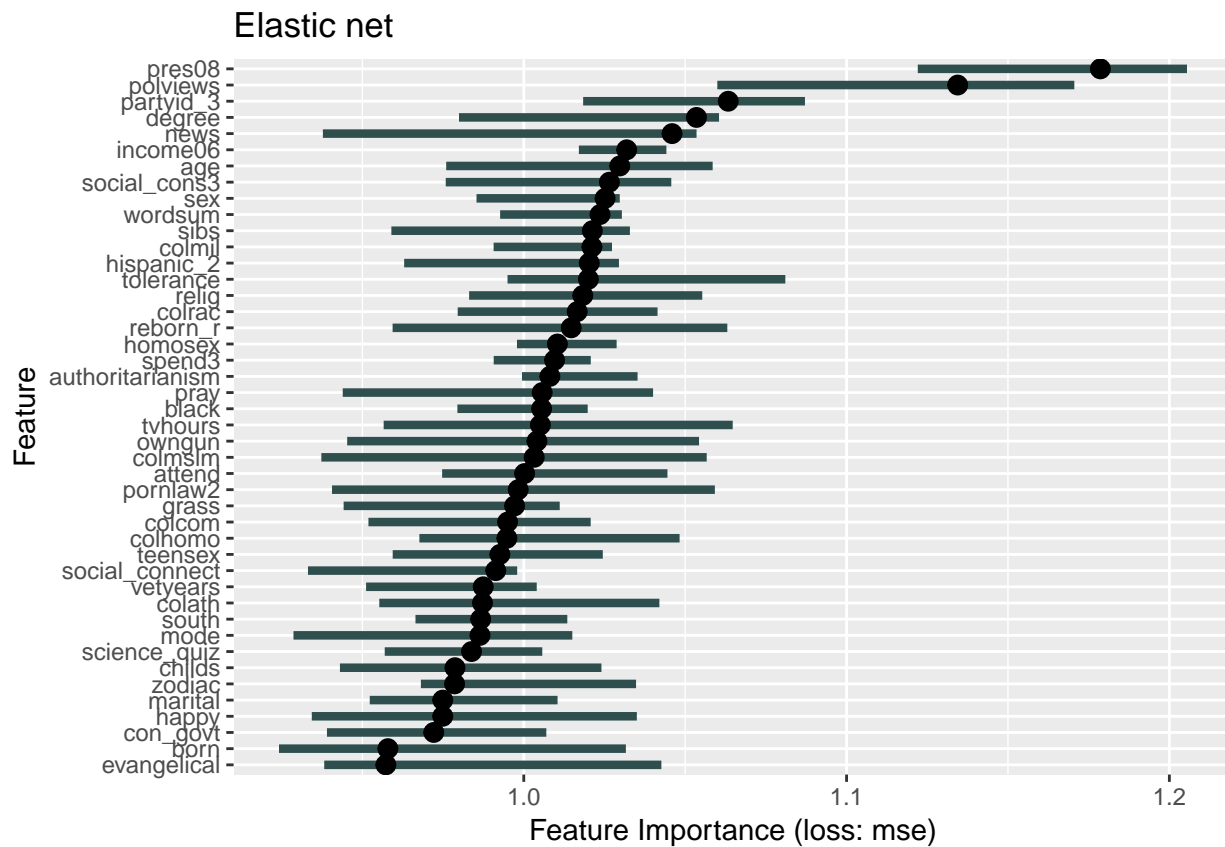
img3



```
head(imp_pcr$results, 5)
```

```
##      feature importance.05 importance importance.95 permutation.error
## 1    pres08    1.0304483    1.066949    1.084274        67.35683
## 2    polviews    1.0184507    1.060965    1.100990        66.97907
## 3     black     0.9948723    1.060885    1.085568        66.97399
## 4 partyid_3     1.0085498    1.050395    1.086067        66.31179
## 5 hispanic_2    0.9863257    1.037460    1.083355        65.49519
```

```
img4
```




```
head(imp_pls$results, 5)
```

##	feature	importance.05	importance	importance.95	permutation.error
## 1	polviews	1.0907496	1.146941	1.206812	65.07771
## 2	pres08	1.0401969	1.104331	1.164965	62.65998
## 3	partyid_3	1.0311443	1.096719	1.117473	62.22808
## 4	hispanic_2	0.9480884	1.038417	1.073760	58.92003
## 5	age	0.9825345	1.033695	1.048706	58.65210

```
img5
```

## MARS



```
head(imp_mars$results, 5)
```

```
##      feature importance.05 importance importance.95 permutation.error
## 1  polviews      1.160049   1.195020      1.219601         66.74976
## 2   pres08      1.107883   1.149817      1.182385         64.22484
## 3 income06      1.029918   1.074666      1.105460         60.02719
## 4     age      1.017997   1.068831      1.078793         59.70123
## 5 partyid_3     1.020129   1.059531      1.066752         59.18180
```

In general, we can find that polviews, pres08 are the most important two features for all of these five model settings; other important features include: partyid\_3, age and income06. I will draw PDP on these variables

```
preds = tibble(name = c("Linear Regression", "PCR", "PLS", "Elastic Net", "MARS"),
  models = list(Linear.Regression = pred_lr,
    Elastic.net = pred_net,
    PCR = pred_pcr,
    PLS = pred_pls,
    MARS = pred_mars
  ))

predictors_pdp <- preds %>%
mutate(
  polviews = map2(models, name, ~ FeatureEffect$new(.x, "polviews", method = "pdp+ice") %>%
    plot() + ggtitle(.y)),
  pres08 = map2(models, name, ~ FeatureEffect$new(.x, "pres08", method = "pdp+ice") %>%
    plot() + ggtitle(.y)),
  partyid_3 = map2(models, name, ~ FeatureEffect$new(.x, "partyid_3", method = "pdp+ice") %>%
```


```

plot() + ggtitle(.y)),
age = map2(models, name, ~ FeatureEffect$new(.x, "age", method = "pdp+ice", center.at = min(gss_train$age)),
plot() + ggtitle(.y)),
inc06 = map2(models, name, ~ FeatureEffect$new(.x, "income06", method = "pdp+ice", center.at = min(gss_train$income06)),
plot() + ggtitle(.y))
)

```


```
predictors_pdp$polviews
```

```
## $Linear.Regression
```




```
##
```

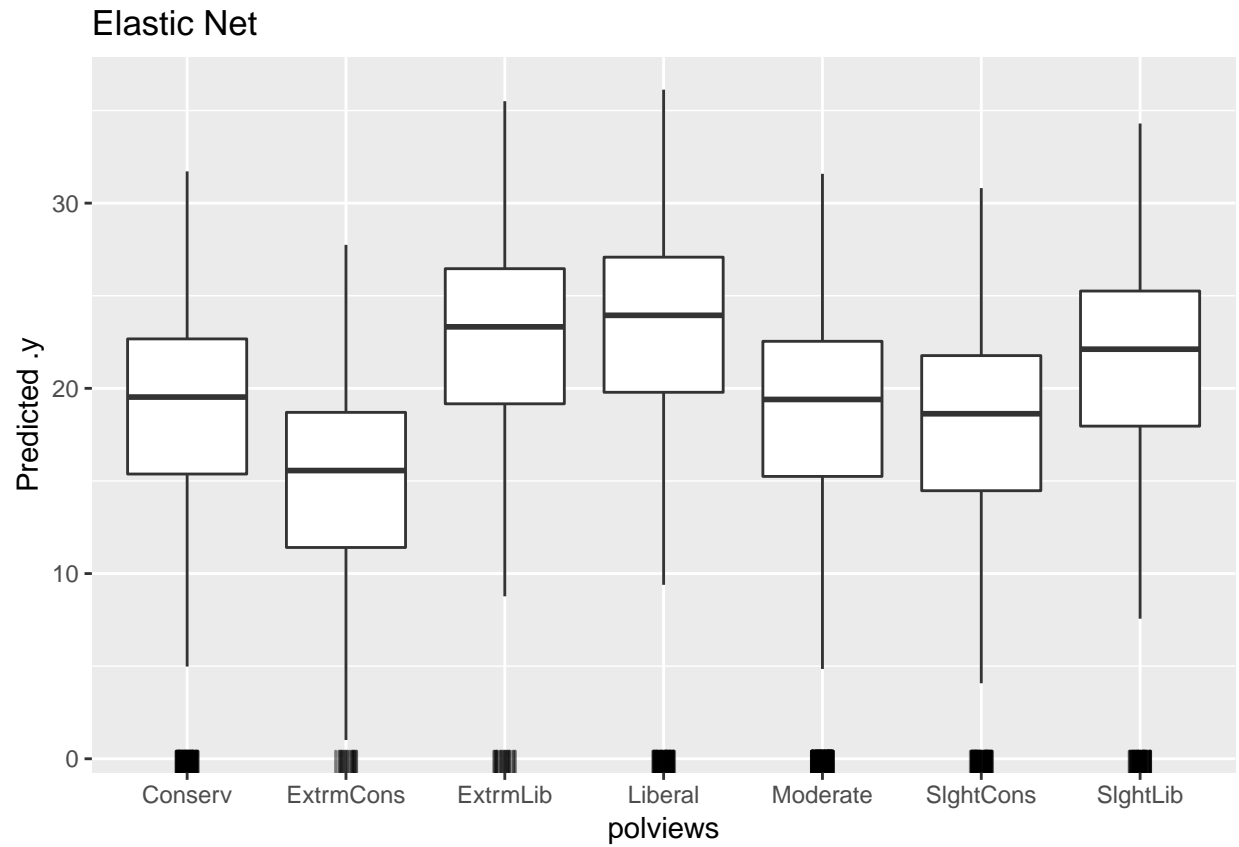
```
## $Elastic.net
```



##  
## \$PCR




##  
## \$PLS



##  
## \$MARS




## MARS




```
predictors_pdp$pres08
```


```
## $Linear.Regression
```



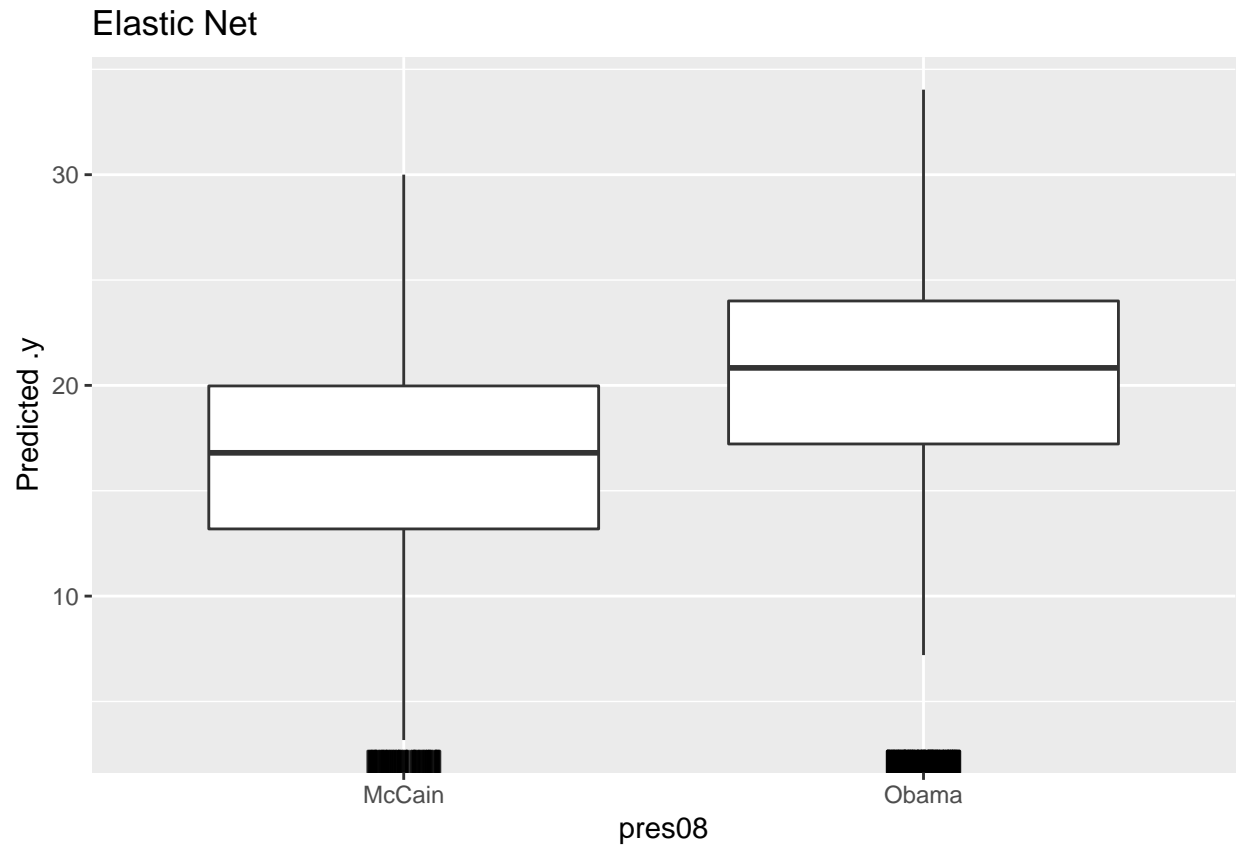
```
##  
## $Elastic.net
```




##  
## \$PCR



##  
## \$PLS

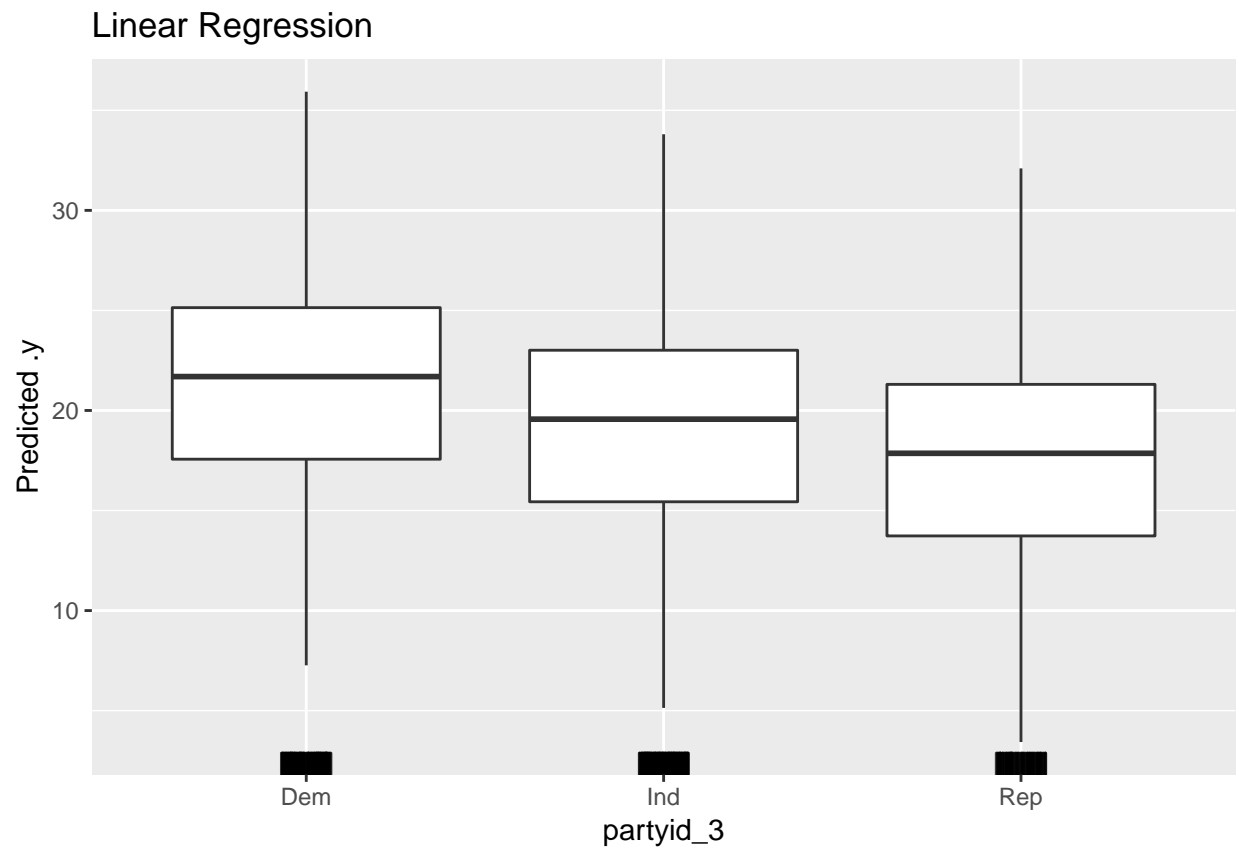


##  
## \$MARS




```
predictors_pdp$partyid_3
```

```
## $Linear.Regression
```




```
##  
## $Elastic.net
```




```
##  
## $PCR
```






##  
## \$PLS

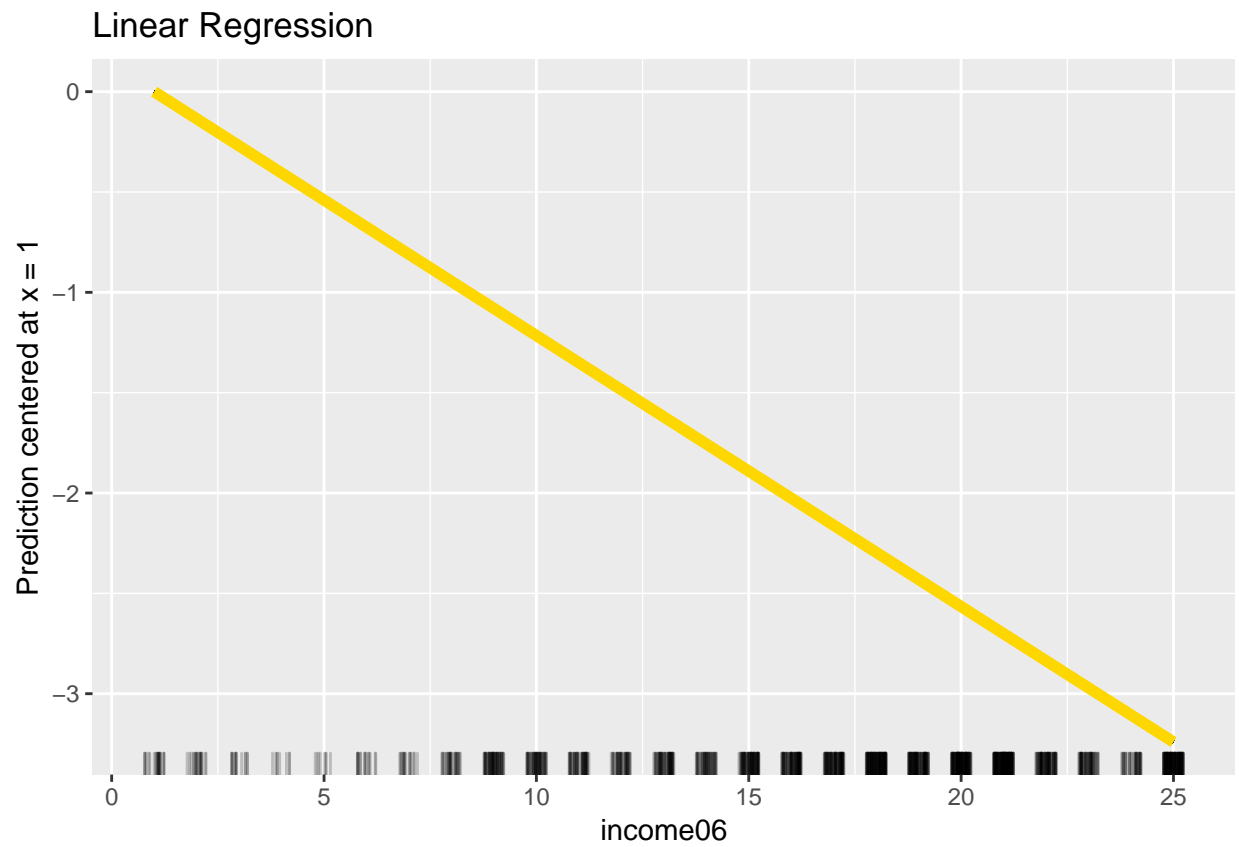


##  
## \$MARS




```
predictors_pdp$inc06
```


```
## $Linear.Regession
```




```
##  
## $Elastic.net
```




##  
## \$PCR



##  
## \$PLS




##  
## \$MARS




```
predictors_pdp$age
```

```
## $Linear.Regression
```







```
##  
## $Elastic.net
```




##  
## \$PCR



##  
## \$PLS




##  
## \$MARS



From these PDPs, we can find that: \* In general, the more liberal the interviewees are, the more egalitarianism they hold. \* Those who voted for Obama are more prone to be egalitarian. \* Democrats favor more egalitarianism. \* In general, the more income people earn, the less egalitarian they are. \* In general, the older people get, the less egalitarian they are.

```
# Linear Regression feature interaction
lr_int = Interaction$new(pred_lr)
lr_int_score = lr_int$results
plot(lr_int) + ggtitle("Linear Regression")
```




```
lr_int_score %>% arrange(-.interaction) %>% head(5)
```

```
##      .feature .interaction
## 1      grass 7.344880e-14
## 2  hispanic_2 7.216728e-14
## 3      colhomo 6.897090e-14
## 4    vetyears 6.669839e-14
## 5       black 6.478122e-14
```

```
# Elastic Net feature interaction
net_int = Interaction$new(pred_net)
net_int_score = net_int$results
plot(net_int) + ggtitle("Elastic Net")
```

## Elastic Net




```
net_int_score %>% arrange(-.interaction) %>% head(5)
```

```
##      .feature .interaction
## 1      colmslm 9.920574e-14
## 2 social_cons3 9.826045e-14
## 3         born 9.528245e-14
## 4         colath 9.329820e-14
## 5      wordsum 9.235250e-14
```

```
# PCR feature interaction
```

```
pcr_int = Interaction$new(pred_pcr)
pcr_int_score = pcr_int$results
plot(pcr_int) + ggtitle("PCR")
```




```
pcr_int_score %>% arrange(-.interaction) %>% head(5)
```

```
##   .feature .interaction
## 1   colath 1.008344e-13
## 2   owngun 8.351261e-14
## 3    happy 7.808475e-14
## 4 vetyears 7.766055e-14
## 5   black 7.547879e-14
```

```
# PLS feature interaction
pls_int = Interaction$new(pred_pls)
pls_int_score = pls_int$results
plot(pls_int) + ggtitle("PLS")
```






```
pls_int_score %>% arrange(-.interaction) %>% head(5)
```

```
##   .feature .interaction
## 1    mode 7.663972e-14
## 2 colhomo 7.234621e-14
## 3    born 6.439649e-14
## 4    black 5.638743e-14
## 5 pornlaw2 5.522438e-14
```

```
# MARS feature interaction
mars_int = Interaction$new(pred_mars)
mars_int_score = mars_int$results
plot(mars_int) + ggtitle("MARS")
```



```
mars_int_score %>% arrange(-.interaction) %>% head(5)
```

```
##      .feature .interaction
## 1      south 1.007076e-13
## 2     colrac 9.173213e-14
## 3      news 8.629674e-14
## 4 authoritarianism 8.516293e-14
## 5     marital 7.838219e-14
```

3. Take the optimal model, apply the test set to the model, and calculate the test set MSE. Does this model generalize well to the test set?

```
# I will go on with the Elastic Net model
predicts = predict(ela.net, gss_test)
y_true = gss_test$egalit_scale
mse = sum((y_true - predicts)^2)/length(y_true)
sqrt(mse)
```

```
## [1] 7.84051
```

Generally, the model generalized well to the test set. In the training process, the average CV RMSE is around 7.72, and it does not inflate very much on the test set.